

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import math
df = pd.read_csv("Titanic_trainn.csv")
df.head(10)
```

Out[1]:

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Ca
0	1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171	7.2500	N
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	0	PC 17599	71.2833	C
2	3	1	3	Heikkinen, Miss. Laina	female	26.0	0	0	STON/O2. 3101282	7.9250	N
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	0	113803	53.1000	C
4	5	0	3	Allen, Mr. William Henry	male	35.0	0	0	373450	8.0500	N
5	6	0	3	Moran, Mr. James	male	NaN	0	0	330877	8.4583	N
6	7	0	1	McCarthy, Mr. Timothy J	male	54.0	0	0	17463	51.8625	I

PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	
7	8	0	3	Palsson, Master. Gosta Leonard	male	2.0	3	1	349909	21.0750	N
8	9	1	3	Johnson, Mrs. Oscar W (Elisabeth Vilhelmina Berg)	female	27.0	0	2	347742	11.1333	N
9	10	1	2	Nasser, Mrs. Nicholas (Adele Achem)	female	14.0	1	0	237736	30.0708	N

In [2]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 12 columns):
#   Column          Non-Null Count  Dtype
---  -
0   PassengerId     891 non-null    int64
1   Survived        891 non-null    int64
2   Pclass          891 non-null    int64
3   Name            891 non-null    object
4   Sex             891 non-null    object
5   Age            714 non-null    float64
6   SibSp          891 non-null    int64
7   Parch          891 non-null    int64
8   Ticket          891 non-null    object
9   Fare           891 non-null    float64
10  Cabin           204 non-null    object
11  Embarked        889 non-null    object
dtypes: float64(2), int64(5), object(5)
memory usage: 83.7+ KB
```

```
In [3]: df.isnull().sum()
```

```
Out[3]: PassengerId      0
Survived      0
Pclass      0
Name      0
Sex      0
Age      177
SibSp      0
Parch      0
Ticket      0
Fare      0
Cabin      687
Embarked      2
dtype: int64
```

```
In [4]: df.describe()
```

```
Out[4]:
```

	PassengerId	Survived	Pclass	Age	SibSp	Parch	Fare
count	891.000000	891.000000	891.000000	714.000000	891.000000	891.000000	891.000000
mean	446.000000	0.383838	2.308642	29.699118	0.523008	0.381594	32.204208
std	257.353842	0.486592	0.836071	14.526497	1.102743	0.806057	49.693429
min	1.000000	0.000000	1.000000	0.420000	0.000000	0.000000	0.000000
25%	223.500000	0.000000	2.000000	20.125000	0.000000	0.000000	7.910400
50%	446.000000	0.000000	3.000000	28.000000	0.000000	0.000000	14.454200
75%	668.500000	1.000000	3.000000	38.000000	1.000000	0.000000	31.000000
max	891.000000	1.000000	3.000000	80.000000	8.000000	6.000000	512.329200

```
In [5]: df.drop("Cabin",axis=1,inplace=True)      # As cabin column contain many
        NAN values so drop it first
df.dropna(inplace=True)
df.head(5)
```

Out[5]:

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	En
0	1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171	7.2500	
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...)	female	38.0	1	0	PC 17599	71.2833	
2	3	1	3	Heikkinen, Miss. Laina	female	26.0	0	0	STON/O2. 3101282	7.9250	
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	0	113803	53.1000	
4	5	0	3	Allen, Mr. William Henry	male	35.0	0	0	373450	8.0500	

In [6]:

```
sex = pd.get_dummies(df["Sex"],drop_first=True) # Create dummies for sex, embark and Pclass column
embark = pd.get_dummies(df["Embarked"],drop_first=True)
Pclass = pd.get_dummies(df["Pclass"],drop_first=True)
#Drop the variables with no use and whose dummies created
df.drop(["Name","Sex","Embarked","Ticket","Pclass"],axis=1,inplace=True)
df.head(5)
```

Out[6]:

	PassengerId	Survived	Age	SibSp	Parch	Fare
0	1	0	22.0	1	0	7.2500
1	2	1	38.0	1	0	71.2833
2	3	1	26.0	0	0	7.9250

	PassengerId	Survived	Age	SibSp	Parch	Fare
3	4	1	35.0	1	0	53.1000
4	5	0	35.0	0	0	8.0500

```
In [7]: df = pd.concat([df,sex,embark,Pclass],axis=1)
df.head(5)
```

Out[7]:

	PassengerId	Survived	Age	SibSp	Parch	Fare	male	Q	S	2	3
0	1	0	22.0	1	0	7.2500	1	0	1	0	1
1	2	1	38.0	1	0	71.2833	0	0	0	0	0
2	3	1	26.0	0	0	7.9250	0	0	1	0	1
3	4	1	35.0	1	0	53.1000	0	0	1	0	0
4	5	0	35.0	0	0	8.0500	1	0	1	0	1

```
In [8]: df.drop("PassengerId",axis=1,inplace=True)
df.head(5)
```

Out[8]:

	Survived	Age	SibSp	Parch	Fare	male	Q	S	2	3
0	0	22.0	1	0	7.2500	1	0	1	0	1
1	1	38.0	1	0	71.2833	0	0	0	0	0
2	1	26.0	0	0	7.9250	0	0	1	0	1
3	1	35.0	1	0	53.1000	0	0	1	0	0
4	0	35.0	0	0	8.0500	1	0	1	0	1

```
In [9]: x = df.drop("Survived",axis=1)
y = df["Survived"]
```

```
In [10]: from sklearn.linear_model import LogisticRegression
```

```
rissh= LogisticRegression()  
rissh.fit(x,y)
```

```
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model\_logistic.py:762: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

```
https://scikit-learn.org/stable/modules/preprocessing.html  
Please also refer to the documentation for alternative solver options:  
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression  
n_iter_i = _check_optimize_result(
```

Out[10]: LogisticRegression()

```
In [12]: predict = rissh.predict(x)  
from sklearn.metrics import classification_report  
classification_report(y,predict)
```

```
Out[12]: '          precision    recall  f1-score   support\n\n         0.81          0.88          0.85         424\n0.70          0.75          0.75          0.75         288\n          0.81          0.81          0.81         712\n\n    macro avg       0.81       0.81       0.81         712\nweighted avg       0.81       0.81       0.81         712'
```

```
In [13]: from sklearn.metrics import confusion_matrix    # Confusion matrix  
confusion_matrix(y,predict)
```

```
Out[13]: array([[373,  51],  
               [ 85, 203]], dtype=int64)
```

```
In [14]: from sklearn.metrics import accuracy_score    # Accuracy of model  
del  
accuracy_score(y,predict)
```

Out[14]: 0.8089887640449438

```
In [15]: print("The intercept is",rissh.intercept_)
```

```
The intercept is [3.8138798]
```

```
In [16]: print("The coeffs of different parameters",rissh.coef_)
```

```
The coeffs of different parameters [[-0.03954496 -0.37442337 -0.0865595  
3  0.00350953 -2.40146613 -0.49214919  
-0.42011897 -0.72821738 -1.97138663]]
```

```
In [ ]:
```